



LOVELY
PROFESSIONAL
UNIVERSITY

Lecture 1 (INT 332) BASICS OF DEVOPS

Revolutionizing Containerization

Why Devops



Understanding the concepts of Virtualization and Containerization



LOVELY
PROFESSIONAL
UNIVERSITY



Imagine a peaceful independent house.




LOVELY
PROFESSIONAL
UNIVERSITY

The Apartment Era (Virtualization)



Infrastructure followed the same path.



One physical server was divided into multiple **virtual machines**.
Each VM felt like its own house — with its own OS, memory, and CPU.

Utilization improved.
Costs reduced.



LOVELY
PROFESSIONAL
UNIVERSITY

Virtualization

- ◎ With the changing times, businesses started looking for solutions to reduce overhead costs, enhance scalability, and standardize application deployment process.
- ◎ They started considering the following two approaches to reduce costs –
 1. Virtualization
 2. Containerization

What is Virtualization?



LOVELY
PROFESSIONAL
UNIVERSITY

- ◎ Virtualization is the process of **partitioning** a physical server into multiple virtual servers.
- ◎ The process of partitioning is carried out using a software called '**hypervisor**'.
- ◎ After partitioning, the virtual servers act and perform just as a physical server.
- ◎ Essentially, it means **using the same hardware** set-up more efficiently, thereby freeing the resources for other tasks or retiring the resources altogether.



Advantages of Virtualization

- ◎ Enhanced performance
- ◎ Promotes agile IT infrastructure
- ◎ Promotes use of resources in optimum manner
- ◎ Better disaster management if any physical hazards take place to resources
- ◎ Better security as the infected VM can be isolated from other VMs and the host server
- ◎ Space saving
- ◎ The capital investment cost on hardware is saved, and thereby maintenance cost is saved, hence overall cost savings for a business.

Drawbacks of Virtualization which lead to concept of containerization



- 1. Heavy Resource Overhead**
- 2. Slow Boot Time and Poor Agility**
- 3. Inefficient Scalability**
- 4. Limited Application Portability**
- 5. Higher Management and Maintenance Complexity**
- 6. Increased Infrastructure Cost**

These drawbacks directly led to the evolution of containerization, which offers lightweight isolation, faster deployment, better portability, and seamless integration with DevOps and microservices architectures.



LOVELY
PROFESSIONAL
UNIVERSITY
PUNJAB (INDIA)

PG Living (Containerization)



High rent.

Fast life.

Limited space.

People started choosing PGs and hostels.

You don't own the building.


You don't manage electricity.

You don't worry about maintenance.

That's exactly what containers do.

- Containers don't carry an entire operating system. They use bare minimum OS.
- They share the host OS and only bring what the application needs.
- They start in seconds. They consume minimal resources.
- You can run dozens — even hundreds — on the same machine.
- Earlier the applications were monolithic , so physical servers were apt,
- later there were components split like backend, frontend, db ,so VM's took care of hosting.
- Now all the applications are turning to microservices, so the containerization tools like docker will serve the purpose.
- Also the portability problem is also solved using the containerization

Need of containers





LOVELY
PROFESSIONAL
UNIVERSITY

Container Runtime

- A container runtime is the software responsible for creating, starting, stopping, and managing containers on a host system.
- It acts as the execution engine that turns a container image into a running container.

In simple words:

“Container runtime is what actually runs containers.”

Need for Container Runtime

- A container image is just a read-only package (application + dependencies).

To run it, the system needs a runtime to:

- Create an isolated environment
- Apply resource limits
- Start the application process
- Monitor and stop it when required

Types of Container Runtimes

1. High-Level Runtimes

Used by users and DevOps tools.

Docker Engine

containerd

CRI-O (Kubernetes-focused)

These handle:

- **Image pulling**
- **Networking**
- **Storage**
- **Container lifecycle**



LOVELY
PROFESSIONAL
UNIVERSITY

2. Low-Level Runtimes

- Responsible for actually running the container process.
- runc (OCI-compliant)

These directly interact with:

- Linux kernel
- Namespaces
- cgroups




LOVELY
PROFESSIONAL
UNIVERSITY

Real-World Analogy

- Container Image → Packed lunch box
- Container Runtime → Person who opens it, serves food, and cleans up

Process Isolation & Namespaces

- Process isolation ensures that each container runs as if it is alone on the system, even though multiple containers share the same OS kernel.
- This isolation is achieved using Linux namespaces.



containers **isolate** a process

Interpretation of Picture

- As shown in this little cat picture, containers are isolated but using networking they still **discover** each other and consume the services if needed.



LOVELY
PROFESSIONAL
UNIVERSITY

Namespaces

- Namespaces are Linux kernel features that partition system resources so that processes see only what belongs to them.
- “Namespaces *create the illusion of a separate system for each container.*”

Namespaces Used in Containers

1. PID Namespace (Process Isolation)

- Each container has its own process tree
- Process IDs inside container start from 1
- A container cannot see host or other container processes

E.g: Inside a container, the main app runs as PID 1, just like init/systemd in a VM.

2. Network Namespace



- **Each container gets:**
 - Its own IP address
 - Its own ports
- **Containers can run apps on the same port without conflict**
- **Example:**
Multiple Nginx containers all listen on port 80 internally.



LOVELY
PROFESSIONAL
UNIVERSITY

3. Mount Namespace

- Each container has its own filesystem view
- File changes inside container do not affect host

4. UTS Namespace

- Allows each container to have its own hostname

5. User Namespace (Advanced)

- Maps container users to non-root users on host
- Improves security

Real-World Analogy behind Namespaces



Namespaces are like separate cabins in a train:

- Same engine (kernel)
- Different passengers
- No visibility between cabins

Control Groups (cgroups) for Resource Limits

Control Groups (cgroups) are Linux kernel features that limit, control, and monitor resource usage of processes.

They ensure that:

- No container can monopolize system resources
- Fair sharing of CPU, memory, and I/O

“Namespaces isolate, cgroups control.”



LOVELY
PROFESSIONAL
UNIVERSITY

Need of cgroups

Without cgroups:

- One container could consume all CPU or RAM
- Host and other containers could crash

Types of Resources Controlled by cgroups

1. CPU

- Limit CPU usage
- Assign CPU shares

2. Memory

- Set maximum memory usage
- Kill container if limit exceeded

3. Disk I/O

- Limit read/write speeds

4. Network (Indirectly)

- Through traffic shaping mechanisms

Real-World Analogy



cgroups are like electricity meters in apartments:

- **Each flat has a limit**
- **One tenant cannot use all power**

Finally for namespaces and cgroups

Containers are lightweight because they do not virtualize hardware.

Instead, a container runtime uses Linux namespaces for process isolation and cgroups for resource control, allowing multiple isolated applications to run efficiently on a shared operating system kernel.



Without cgroups

(Problem Scenario)

HOST SYSTEM

CPU: 100% Memory: 16 GB

Container A (No Limits)

Uses **90%** CPU

Uses **14 GB** RAM

Container B
CRASHED

Container C
CRASHED

With cgroups

(Controlled)

HOST SYSTEM

CPU: 100% Memory: 16 GB

Container A (cgroups)

CPU Limit : **30%**

RAM Limit: **4 GB**

Container B

CPU: **30%** RAM: **6 GB**

Container C

CPU: **40%** RAM: **6 GB**

Kernel Enforces Limits



cgroups
Resource Limits



Namespaces
Isolation

No Container Can Monopolize Resources



LOVELY
PROFESSIONAL
UNIVERSITY

Today's Task

Task 1: A company uses containers for microservices. One faulty container crashes the host due to high memory usage.

Questions:

- **Which container feature was misconfigured or missing?**
- **Which kernel mechanism should have prevented this?**
- **Would namespaces alone solve this issue? Why?**



Task 2: A server is running three containers:

- Container A runs a CPU-intensive task
- Container B runs a web server
- Container C runs a database

Questions:

- Which Linux feature prevents Container A from consuming all CPU?
- Which feature ensures Container B cannot see processes of Container C?
- Which namespace allows all three containers to use port 80 internally?

Container Image

A container image is:

- A **read-only blueprint**
- **Contains:**
 - Base OS (minimal)
 - Application code
 - Libraries & dependencies
 - Runtime
 - Configuration

Images are immutable – once built, they do not change.



LOVELY
PROFESSIONAL
UNIVERSITY

Example

python:3.11-slim is the image

This image already includes:

- Linux OS
- Python 3.11
- Standard Python libraries


Each image is built as a stack of layers.

Characteristics of Layers

- **Each instruction in a Dockerfile = one layer**
- **Layers are:**
 - **Immutable**
 - **Cached**
 - **Reused across images**



Image Layers



Benefits of Layered Images



LOVELY
PROFESSIONAL
UNIVERSITY

Faster Builds

- Docker reuses unchanged layers (cache)

Smaller Storage

- Same layers shared between images

Faster Deployment

- Only missing layers are downloaded

Container Image vs Container




Image Registry



An image registry is a centralized repository used to store, manage, version, and distribute container images.

It enables:

- **Sharing images across teams and environments**
- **Consistent deployment in development, testing, and production**
- **Integration with CI/CD pipelines**

In DevOps, an image registry plays the same role as:

- **GitHub → for source code**
- **Maven Central → for Java libraries**

Why Image Registries are Required



Without a registry:

- Each system must build images independently
- High inconsistency across environments
- Manual software installation
- No version control of application environments

With a registry:

- Build once, deploy everywhere
- Faster deployments
- Version-controlled environments
- Rollback capability
-

Types of Image Registries



3.1 Public Image Registries

Public registries allow open access to images.

Characteristics:

- **Images are publicly visible**
- **Anyone can pull images**
- **Often free**
- **Used for base images and open-source applications**

Private Image Registries



Private registries restrict access using authentication and authorization.

Characteristics:

- **Secure and controlled access**
- **Used in enterprises**
- **Stores proprietary applications**
- **Integrated with IAM and RBAC**

Examples:

- **AWS Elastic Container Registry (ECR)**
- **Azure Container Registry (ACR)**
- **GitHub Container Registry (GHCR)**
- **On-premise private registry**

Image Naming Convention

Images follow a standard naming

format:<registry>/<namespace>/<image>:<tag>

Example:`docker.io/harpreet/webapp:v1`

Component	Meaning
registry	Location of image
namespace	User or organization
image	Application name
tag	Version or label



Tags and Versioning

Tags identify different versions of the same image.

Common Tagging Strategies

Tag	Purpose
latest	Default (not recommended in production)
v1, v2	Versioned releases
dev	Development build
qa	Testing build
prod	Stable production build

Image Distribution Process

Image distribution refers to how images move from developers to production systems.

Distribution Flow:

- Developer builds an image locally
- Image is pushed to a registry
- Registry stores the image securely
- Servers pull the image when required
- Containers are created from the image

Push and Pull Operations

Push Operation

- **Uploads image layers from local system to registry**
- **Requires authentication**
- **Only new layers are pushed**

Pull Operation

- **Downloads image layers from registry**
- **Uses caching to avoid re-downloading existing layers**
- **Faster on repeated deployments**

Role of Image Registries in CI/CD Pipelines

In modern DevOps pipelines:

- Code is committed to Git
- CI system builds the container image
- Image is pushed to registry
- CD system pulls the image
- Application is deployed automatically

In Actual, Registry acts as the bridge between CI and CD.



LOVELY
PROFESSIONAL
UNIVERSITY

Comparison Between Public and Private Registry

Feature	Public Registry	Private Registry
Access	Open	Restricted
Security	Lower	High
Cost	Mostly free	Paid
Use Case	Learning, open source	Enterprise, production

QUIZ

1.What is the primary purpose of an image registry?

- A. To execute containers
- B. To store and distribute container images
- C. To build container images
- D. To monitor container performance

2. Which component of an image name represents the version of the image?

- A. Registry
- B. Namespace
- C. Image name
- D. Tag



LOVELY
PROFESSIONAL
UNIVERSITY

3. Which type of image registry is typically used in enterprise environments?

- A. Public registry
- B. Open registry
- C. Private registry
- D. Anonymous registry

4. In CI/CD pipelines, the image registry mainly acts as:

- A. Source code repository
- B. Testing environment
- C. Bridge between CI and CD
- D. Container runtime



LOVELY
PROFESSIONAL
UNIVERSITY

Introduction to Docker Hub

Centralized repository for Docker images

Features:

- **Public and private repositories**
- **Official images**
- **Automated builds**
- **Importance in DevOps**



LOVELY
PROFESSIONAL
UNIVERSITY


Docker Architecture

Key Components:

- - Docker Client
- - Docker Daemon
- - Docker Images
- - Containers
- - Docker Registry (Docker Hub)

Docker Architecture

Docker Architecture Overview



Analogy:



Example Workflow:





LOVELY
PROFESSIONAL
UNIVERSITY

Docker Life Cycle

Stages:

- - Build: Create an image
- - Ship: Share images via Docker Hub
- - Run: Start containers

Container states:

- - Created
- - Running
- - Paused
- - Stopped

Why is Docker faster than VM?

- Containers share the Host OS Kernel
- No Guest OS boot needed
- Less resource usage & faster startup

Docker Containers




 Shared Host OS Kernel

Docker Containers

 Fast Startup |  Low Resources

Virtual Machines



 Separate Guest OS + Hypervisor

Virtual Machines

 Slow Startup |  High Resources



LOVELY
PROFESSIONAL
UNIVERSITY

Conclusion

Key takeaways:

- Docker simplifies software deployment
- Docker Hub enhances collaboration
- Docker's architecture and lifecycle enable efficient workflows
- Hyper-V integration broadens Docker's capabilities